

---

---

# Prediction on the rating of fine food and evaluation for the helpfulness of review based on the Amazon Fine Food Review

Capstone Project by Liyan Cen

---

# The problem & who might care?

Customer reviews is an important part of the Amazon shopping experience. It is so far the major reference and indicator for Amazon to evaluate the quality of their products, and the performance of their sellers.


Additionally, it is a great tool for customers to compare between products. Amazon always tend to place the most helpful reviews at the first place of the list as a reference for customers, and also as a standard for reviewers to follow. The company carefully considers any changes to the ratings and review system to ensure customers and sellers continue to trust it.

---

# Data Information

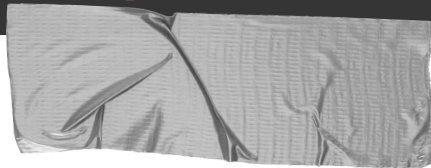
**This dataset consists of reviews of fine foods from amazon. The data span a period of more than 10 years, including all ~500,000 reviews up to October 2012. Reviews include product and user information, ratings, and a plain text review.**

# Data Cleaning




- This dataset consists of reviews of fine foods from amazon. The data span a period of more than 10 years, including all ~500,000 reviews up to October 2012. Reviews include product and user information, ratings, and a plain text review.
- Because there is memory limitation of our CPU, we first cross out the “HelpfulnessDenominator” value that is smaller than 10. Also, we cross out all rows which has null values, and make sure there is no duplicate contents. As a result, the length of data changes from 568,454 to 21,463.

# Data Cleaning - Continued



- In order to make prediction models with better performance along with less noises or outliers, we first reprocess the text to make our feature set more correlated with the target variable. By doing that, we create a function called “normalize\_corpus” to remove punctuations, digits, html tags, accented characters, special characters, and stop words. At the same time, we lower all strings of each text.

# Data Cleaning - Continued



- For both the classification and regression models, we use text as our feature set by converting each word from each text into numerators by “TfidfVectorizer”. Frequent words across all documents may tend to overshadow other terms in the feature set. The TF-IDF model is functioned to solve this issue by using a scaling or normalizing factor in its computation.

## Data Wrangling and Cleaning - “Ratings vs. Reviews”

For this section, we use text feature as our feature set, and score as our target variable. Classification models are used because our model predicts and classifies reviews into a score ranging from 1 to 5.


## Data Wrangling and Cleaning - “Helpfulness vs. Reviews”

For this section, we use text feature as our feature set, and “Helpful” as our target variable. We get “Helpfulness ratio” by having “HelpfulnessNumerator” divided by “HelpfulnessDenominator”. Classification models are used because our model predicts and classifies reviews into either “0” or “1”, while ratio greater than the mean will be equal to “1”, and less than the mean will be equal to “0”.



# Data Exploration

Based on the graph, we can see that around 70% of the text features appeared in about 4 or 5 documents, so we set our `min_df` to be “5”, which means we cut off the words that occur in 5 documents or less.



# Modeling

We use supervised learning algorithms to build predictive models, and find out the best model based on their performance, specifically regarding to each Accuracy and Precision on testing sets. As we mentioned earlier, we will have the cleaned and vectorized “reviews” as our feature set, and the “ score” and “Helpful” as our predictors. We perform Logistics regression, Naives Bayes, LinearSVC, Random Forests, Gradient Boosting and SGDClassifier as our models, and compare accuracy for each to figure out the best model, and parameter.

# Ratings vs. Reviews - Naive Bayes

By referring to the data above, we can tell that the model does not perform very well, this might be caused by the oversimplified assumption.

Accuracy on training set: 0.6767744085304899

Accuracy on testing set: 0.6559390547263682

## Model Performance metrics:

-----  
Accuracy: 0.6559  
Precision: 0.7426  
Recall: 0.6559  
F1 Score: 0.5752

## Model Classification report:

-----  
                  precision    recall    f1-score    support  
                  1          0.76      0.63      0.69      1873  
                  2          0.92      0.05      0.09      474  
                  3          1.00      0.00      0.00      480  
                  4          1.00      0.03      0.07      534  
                  5          0.62      0.97      0.76      3071  
avg / total          0.74      0.66      0.58      6432

## Prediction Confusion Matrix:

-----  
                  Predicted:  
                          1      2      3      4      5  
Actual: 1          1189      0      0      0      684  
          2          138     22      0      0      314  
          3          115      0      1      0      364  
          4          39      0      0     18      477  
          5          80      2      0      0     2989

# Ratings vs. Reviews - Logistic Regression

By referring to the data above, we can tell that the model performs pretty well, especially for score 1 and score 5.

Comparatively, score 2,3, and 4 do not perform as well as the other ones; this might be caused by the fact we do not have as many reviews with score 2 to 4 as we have for score 1 and 5 from our dataset.

Accuracy on training set: 0.87350883038987

Accuracy on testing set: 0.7184390547263682

Model Performance metrics:

-----  
Accuracy: 0.7184  
Precision: 0.7222  
Recall: 0.7184  
F1 Score: 0.7199

Model Classification report:

-----  
                  precision    recall  f1-score    support  
                  -----  
                  1          0.75      0.79      0.77      1815  
                  2          0.44      0.43      0.43      502  
                  3          0.43      0.43      0.43      494  
                  4          0.40      0.42      0.41      524  
                  5          0.86      0.82      0.84      3097  
                  -----  
avg / total          0.72      0.72      0.72      6432

Prediction Confusion Matrix:

-----  
                  Predicted:  
                          1      2      3      4      5  
Actual: 1          1432  117  88  64  114  
          2          144  215  55  31  57  
          3          102  55  214  52  71  
          4          39  31  51  220  183  
          5          203  73  93  188  2540

# Ratings vs. Reviews - LinearSVC

Given the fact that this project is to predict on text classification, the feature sets and the target variables may not be linearly related. Also, due to the large amount of unique words in the dataset, 44,755 features in specific, we choose to perform SVC models, one of models in SVM, with a non-linear kernel, RBF.

By referring to the data above, we can tell that the LinearSVC model performs pretty well, and actually a little bit better than logistic regression.

Accuracy on training set: 0.9538820393202266

Accuracy on testing set: 0.7580845771144279

Model Performance metrics:

-----  
Accuracy: 0.7581  
Precision: 0.7424  
Recall: 0.7581  
F1 Score: 0.7424

Model Classification report:

-----  
                    precision    recall    f1-score    support  
1          0.74          0.83          0.79        1815  
2          0.64          0.38          0.48         502  
3          0.61          0.40          0.48         494  
4          0.56          0.36          0.44         524  
5          0.81          0.90          0.85        3097  
avg / total          0.74          0.76          0.74        6432

Prediction Confusion Matrix:

-----  
                    Predicted:  
                          1      2      3      4      5  
Actual: 1          1508    48    39    24    196  
         2          172   193    32    19     86  
         3          117    24   196    32   125  
         4           50    17    17   189   251  
         5          179    18    37    73   2790

# Ratings vs. Reviews - SGDClassifier

Accuracy on training set: 0.8421859380206598

Accuracy on testing set: 0.755907960199005

While Logistics Regression and Support Vector Machines have a disadvantage of discriminating linear classifiers under convex loss functions, we try out SGDClassifier model to compensate this issue.

Based on the Accuracy and Precision on the testing set, this model performs just as well as LinearSVC.

## Model Performance metrics:

-----  
Accuracy: 0.7559  
Precision: 0.7512  
Recall: 0.7559  
F1 Score: 0.7207

## Model Classification report:

-----  
                  precision    recall    f1-score    support  
                  -----  
                  1          0.74      0.84      0.79      1791  
                  2          0.74      0.29      0.42      456  
                  3          0.82      0.26      0.39      468  
                  4          0.64      0.20      0.30      536  
                  5          0.77      0.94      0.85      3181  
  
avg / total          0.75      0.76      0.72      6432

## Prediction Confusion Matrix:

-----  
                  Predicted:  
                          1      2      3      4      5  
Actual: 1          1508    14      7      4     258  
          2          194   132    10    12    108  
          3          124    13   120    23    188  
          4           63    11      5   106    351  
          5          152      9      4    20   2996

# Ratings vs. Reviews - Decision Tree

Based on the “Accuracy on testing set” above, we can tell that the Decision Tree model does not perform as well as the other ones. This might be due to the nature of decision tree model, which it tends to make the most optimal decision at each step, instead of taking into account the global optimum. The other reason may also be that decision tree trains data based on the results of the last training data.

Accuracy on training set: 1.0

Accuracy on testing set: 0.6731965174129353

## Model Performance metrics:

-----  
Accuracy: 0.6732  
Precision: 0.6614  
Recall: 0.6732  
F1 Score: 0.6658

## Model Classification report:

-----  
                  precision      recall   f1-score     support  
          1         0.68         0.70         0.69         1815  
          2         0.47         0.39         0.43         502  
          3         0.47         0.38         0.42         494  
          4         0.44         0.36         0.40         524  
          5         0.75         0.80         0.78         3097  
avg / total         0.66         0.67         0.67         6432

## Prediction Confusion Matrix:

-----  
                  Predicted:  
                          1      2      3      4      5  
Actual: 1         1265     88     62     45     355  
         2         120     195     35     19     133  
         3         103     36     186     31     138  
         4          80     15     30     191     208  
         5         289     79     84     152     2493

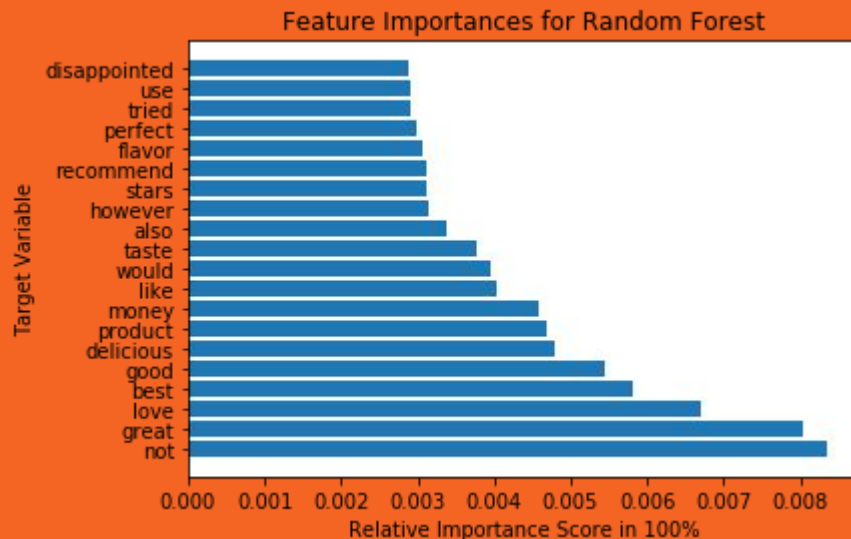
# Ratings vs. Reviews - Random Forests

Based on the top 20 feature importances listed above, we get a sense of which word variables have the most effect in our models. As we can tell, the model performs pretty well because all the words listed above are the common words we see and use in reviews.

Also, by referring to the accuracy, we can tell the this model performs pretty well, and is actually the third best model of all besides LinearSVC and SGDClassifier. Additionally, score 2,3, and 4 perform well in this model.

**Accuracy on training set: 0.9910696434521826**

**Accuracy on testing set: 0.7391169154228856**






# Ratings vs. Reviews - Gradient Boosting

Even though Random Forests performs well on this dataset, it does have its limitations like the features were chosen randomly with replacement. Hence, we try out another popular decision tree model, Gradient Boosting, to fix this issue by training on data instances that had been modeled poorly in the overall system before.

According to the output, Gradient Boosting does not perform as well as Random Forests. This might be due to the fact that Gradient Boosting' training is based on last training result whereas Random Forests is trained independently from the rest.

Accuracy on training set: 0.738487170943019  
Accuracy on testing set: 0.6906094527363185



# Ratings vs. Reviews - Hyperparameter tuning for the best parameter

Based on all the models listed above, it turns out that the Random Forest, SGDClassifier, and LinearSVC give us the best results. In order to re-confirm which model is the best for our dataset, we use grid search to figure out the best hyperparameter by performing each model one more time.

By referring to the output for each model, it turns out that the Random Forest gives us the best parameter.

Accuracy on training set: 1.0

Accuracy on testing set: 0.7657027363184079

Model Performance metrics:

-----  
Accuracy: 0.7657  
Precision: 0.811  
Recall: 0.7657  
F1 Score: 0.7385

Model Classification report:

-----  
                  precision    recall    f1-score    support  
1          0.81        0.78        0.79        1873  
2          1.00        0.35        0.52         474  
3          1.00        0.31        0.48         480  
4          1.00        0.30        0.46         534  
5          0.72        0.97        0.83        3071  
  
avg / total         0.81        0.77        0.74        6432

Prediction Confusion Matrix:

-----  
                  Predicted:  
                          1      2      3      4      5  
Actual: 1         1456      0      0      0     417  
         2         120    165      0      0     189  
         3         102      0    151      0     227  
         4          37      0      0    161     336  
         5          79      0      0      0    2992

# Helpful vs. Reviews

## Naive Bayes

Accuracy on training set: 0.7872736954206603

Accuracy on testing set: 0.7605218201584097

Model Performance metrics:

-----  
Accuracy: 0.7605

Precision: 0.78

Recall: 0.7605

F1 Score: 0.7184

Model Classification report:

-----  
                  precision    recall    f1-score    support  
                  0         0.85      0.30      0.44      2032  
                  1         0.75      0.98      0.85      4407  
avg / total         0.78      0.76      0.72      6439

Prediction Confusion Matrix:

-----  
                  Predicted:  
                          0      1  
Actual: 0          600  1432  
          1         110  4297

## Logistic Regression

Accuracy on training set: 0.86395101171459

Accuracy on testing set: 0.8024537971734741

Model Performance metrics:

-----  
Accuracy: 0.8025

Precision: 0.8148

Recall: 0.8025

F1 Score: 0.8062

Model Classification report:

-----  
                  precision    recall    f1-score    support  
                  0         0.66      0.77      0.71      2032  
                  1         0.89      0.82      0.85      4407  
avg / total         0.81      0.80      0.81      6439

Prediction Confusion Matrix:

-----  
                  Predicted:  
                          0      1  
Actual: 0         1571  461  
          1         811  3596

# Helpful vs. Reviews

## LinearSVC

Accuracy on training set: 0.9959398296059638

Accuracy on testing set: 0.8141015685665476

Model Performance metrics:

-----  
Accuracy: 0.8141  
Precision: 0.8152  
Recall: 0.8141  
F1 Score: 0.8146

Model Classification report:

-----  
                  precision    recall  f1-score   support  
                  0          0.70     0.72     0.71     2032  
                  1          0.87     0.86     0.86     4407  
avg / total         0.82      0.81     0.81     6439

Prediction Confusion Matrix:

-----  
                  Predicted:  
                          0      1  
Actual: 0          1455   577  
          1          620  3787

## SGDClassifier

Accuracy on training set: 0.8893104366347178

Accuracy on testing set: 0.8276129833825128

Model Performance metrics:

-----  
Accuracy: 0.8276  
Precision: 0.8239  
Recall: 0.8276  
F1 Score: 0.8218

Model Classification report:

-----  
                  precision    recall  f1-score   support  
                  0          0.78     0.63     0.70     2032  
                  1          0.84     0.92     0.88     4407  
avg / total         0.82      0.83     0.82     6439

Prediction Confusion Matrix:

-----  
                  Predicted:  
                          0      1  
Actual: 0          1273   759  
          1          351  4056

# Helpful vs. Reviews

## Decision Tree

Tree One, Accuracy on training set: 0.9987353567

Tree One, Accuracy on testing set: 0.7808665941

Tree Two, Accuracy on training set: 0.68450479

Tree Two, on testing set: 0.6903245845628203

Model Performance metrics:

-----  
Accuracy: 0.7809  
Precision: 0.7812  
Recall: 0.7809  
F1 Score: 0.781

Model Classification report:

-----  
                  precision    recall  f1-score  support  
                  0          0.65     0.66     0.65     2032  
                  1          0.84     0.84     0.84     4407  
avg / total         0.78     0.78     0.78     6439

Prediction Confusion Matrix:

-----  
                  Predicted:  
                          0      1  
Actual: 0         1332   700  
         1          711  3696

Model Performance metrics:

-----  
Accuracy: 0.6903  
Precision: 0.713  
Recall: 0.6903  
F1 Score: 0.5743

Model Classification report:

-----  
                  precision    recall  f1-score  support  
                  0          0.76     0.03     0.05     2032  
                  1          0.69     1.00     0.81     4407  
avg / total         0.71     0.69     0.57     6439

Prediction Confusion Matrix:


-----  
                  Predicted:  
                          0      1  
Actual: 0          55   1977  
         1          17   4390

# Helpful vs. Reviews

## Gradient Boosting

Accuracy on training set: 0.7279685835995741


Accuracy on testing set: 0.7261997204534866



## Random Forests

Accuracy on training set: 0.9932774227902024

Accuracy on testing set: 0.8120826215250815



# Helpful vs. Reviews - Hyperparameter tuning for the best parameter

Based on all the models listed above, it turns out that the Random Forest, SGDClassifier, Logistic Regression and LinearSVC give us the best results. In order to re-confirm which model is the best for our dataset, we use grid search to figure out the best hyperparameter by performing each model one more time.

By referring to the output for each model, it turns out that the Random Forest gives us the best parameter.

Accuracy on training set: 0.9987353567625133

Accuracy on testing set: 0.8350675570740799

Model Performance metrics:

-----  
Accuracy: 0.8351  
Precision: 0.8375  
Recall: 0.8351  
F1 Score: 0.8248

Model Classification report:

-----  
                    precision      recall   f1-score     support  
                    0          0.85      0.58      0.69      2032  
                    1          0.83      0.95      0.89      4407  
avg / total          0.84      0.84      0.82      6439

Prediction Confusion Matrix:

-----  
                    Predicted:  
                                0          1  
Actual: 0          1171      861  
          1          201     4206

## Limitations

We lack the knowledge of the major market for each category of fine food. In other words, the reviewers are not narrowed down in ages, jobs, or other backgrounds. Hence, the degree of helpfulness could be somewhat biased.

While positive reviews are helping Amazon and sellers generating more profits, fake reviews could be a concern that hurt their reputations. For example, a Washington Post examination found that for some popular product categories, such as Bluetooth headphones and speakers, the vast majority of reviews appear to violate Amazon's prohibition on paid reviews. Such reviews have certain characteristics, such as repetitive wording that people probably cut and paste in. Regarding to this issue, we lack the knowledge of whether the review is from a verified user.



---

## Conclusions

For both sections, the RandomForestRegressor gives us the best  $R^2$  value among all the models we have performed. Also, it give us the best hyperparameter and model.